

Comparing Pretrained Transformer Models on the GLUE SST-2 Sentiment Classification Task

Ye Cai

Department of Statistics
University of Michigan
Email: cyea@umich.edu

Abstract—Pretrained Transformer language models have become the standard backbone for many natural language processing tasks. In this project we fine-tune and compare four widely used pretrained encoders—DistilBERT, BERT-base-uncased, ALBERT-base-v2, and RoBERTa-base—on the SST-2 sentiment classification task from the GLUE benchmark. Using the Hugging Face transformers and datasets libraries in a unified Python pipeline, we keep preprocessing, optimization hyperparameters, and the evaluation protocol fixed so that performance differences can be attributed mainly to the model architectures. We report validation accuracy, training runtime, and throughput on a single T4 GPU, and analyze the training and evaluation curves logged with Weights and Biases. RoBERTa-base attains the best validation accuracy (around 93.3%), while DistilBERT trains substantially faster at the cost of several points of accuracy. BERT-base and ALBERT-base-v2 lie between these extremes. Our results quantify the accuracy–efficiency trade-off among these popular models and provide practical guidance for choosing a pretrained encoder for resource-constrained sentiment classification applications.

Index Terms—Transformers, sentiment analysis, text classification, GLUE, SST-2, Hugging Face

I. INTRODUCTION

Sentiment classification aims to automatically determine whether a piece of text expresses a positive or negative opinion. It is a core building block for many downstream applications, including product review mining, social media monitoring, and user feedback analysis. In recent years, large pretrained Transformer language models have dramatically improved the state of the art on sentiment classification and other natural language processing (NLP) tasks by providing powerful contextual representations that can be adapted to new tasks with relatively little labeled data.

Among the most influential models are BERT, RoBERTa, ALBERT, and DistilBERT. BERT-base [1] introduced bidirectional Transformer encoders pretrained with masked language modeling and next-sentence prediction, and quickly became a strong baseline on the GLUE benchmark. RoBERTa [2] later showed that with more data and a better training recipe, the same architecture can reach even stronger performance. ALBERT [3] reduces memory footprint and improves parameter efficiency by factorizing embeddings and sharing weights across layers. DistilBERT [4] applies knowledge distillation to obtain a smaller and faster model while retaining most of BERT’s accuracy. All of these models are widely available

through the Hugging Face transformers library and are commonly used as backbones for text classification.

While benchmark leaderboards often report the best achievable accuracy for each model, practitioners frequently care about the trade-off between performance and efficiency under realistic resource constraints. For a given downstream task, a heavier model may yield slightly better accuracy but require much longer training time and higher inference cost. Conversely, a lighter model may be attractive for deployment but could sacrifice several points of accuracy. Understanding these trade-offs on a concrete task is important for making informed model choices in practice.

In this project, we focus on the SST-2 sentiment classification task from the GLUE benchmark and ask the following questions:

- How do four popular pretrained encoders—DistilBERT, BERT-base-uncased, ALBERT-base-v2, and RoBERTa-base—compare in terms of validation accuracy when finetuned under a shared training pipeline?
- How do their training runtimes and throughput on a single T4 GPU differ, and what efficiency gains do the smaller models actually offer?
- What practical accuracy–efficiency trade-off emerges from these results, and how might it guide model selection for resource-constrained sentiment classification applications?

To answer these questions, we build a unified fine-tuning pipeline in Python using the Hugging Face transformers and datasets libraries. We keep the data preprocessing and optimization hyperparameters fixed across models so that differences in performance can be attributed primarily to the model architectures. We then analyze both the final validation metrics and the training dynamics logged with Weights and Biases to obtain a detailed comparison of the four models.

II. METHOD

A. Problem Formulation

We consider binary sentiment classification on single sentences. Let x denote an input sentence and $y \in \{0, 1\}$ denote its sentiment label, where 0 corresponds to negative and 1 to positive sentiment. Given a labeled training set $\{(x_i, y_i)\}_{i=1}^N$, our goal is to learn a classifier $f_\theta(x)$ parameterized by θ that

estimates the conditional probability $p_\theta(y | x)$ and predicts the most likely label

$$\hat{y} = \arg \max_{y \in \{0,1\}} p_\theta(y | x).$$

We model f_θ as a pretrained Transformer encoder followed by a task-specific classification head, and we fine-tune all parameters on the SST-2 training set.

B. Dataset

We use the Stanford Sentiment Treebank 2 (SST-2) [6] as provided through the GLUE benchmark [5]. SST-2 contains movie review sentences labeled as expressing either positive or negative sentiment. Following the GLUE setup, we use the official train and validation splits and treat the task as binary single-sentence classification. After filtering out examples with neutral sentiment, the training set contains approximately 67k sentences and the validation set contains about 0.9k sentences.

We load the dataset using the Hugging Face datasets library:

```
from datasets import load_dataset
dataset = load_dataset("glue", "sst2")
```

The resulting `DatasetDict` provides separate train and validation splits that we use throughout our experiments.

C. Models

We compare four pretrained Transformer encoders, all accessed via the Hugging Face transformers library:

- **DistilBERT-base-uncased** [4]: a distilled version of BERT-base with approximately 66M parameters and 6 Transformer layers, designed to be smaller and faster while retaining most of BERT’s performance.
- **BERT-base-uncased** [1]: the original 12-layer Transformer encoder with 110M parameters, pretrained with masked language modeling and next-sentence prediction.
- **ALBERT-base-v2** [3]: a parameter-efficient variant of BERT that factorizes the embedding matrix and shares weights across layers, reducing memory usage while maintaining competitive accuracy.
- **RoBERTa-base** [2]: an optimized version of BERT trained on more data with a modified pretraining objective and larger batch sizes, often achieving stronger performance on GLUE.

For each model, we use the corresponding AutoTokenizer and AutoModelForSequenceClassification classes. The classification head is a randomly initialized linear layer mapping the pooled representation to two output logits.

D. Training Setup

1) *Preprocessing*: We tokenize each sentence using the model-specific tokenizer with WordPiece or byte-pair encoding, truncate sequences longer than 128 tokens, and apply padding so that all examples in a mini-batch have the same length. The preprocessing function is applied to the entire dataset using the `map` method from the datasets library, and the results are cached to avoid repeated computation across runs.

...	model	eval_accuracy	train_runtime
0	albert-base-v2	0.924312	2501.4015
1	bert-base-uncased	0.923165	2702.7464
2	distilbert-base-uncased	0.909406	1741.1033
3	roberta-base	0.933486	2964.0291

Fig. 1. Validation accuracy of four pretrained models on the SST-2 dev set.

TABLE I
VALIDATION ACCURACY AND TRAINING RUNTIME ON SST-2.

Model	Dev accuracy	Train runtime (s)
ALBERT-base-v2	0.9243	2501.4
BERT-base-uncased	0.9232	2702.7
DistilBERT-base	0.9094	1741.1
RoBERTa-base	0.9335	2964.0

2) *Optimization*: To enable a fair comparison, we adopt a shared set of fine-tuning hyperparameters for all four models. Unless otherwise noted, we train for 5 epochs with a batch size of 16, using the AdamW optimizer with a learning rate of 2×10^{-5} and weight decay of 0.01. We employ a linear learning rate schedule with warmup over 10% of the total training steps. Evaluation on the validation set is performed at the end of each epoch, and we select the checkpoint with the best validation metric as the final model.

We implement training using the Trainer API from Hugging Face transformers, which handles the training loop, gradient accumulation, and evaluation. Accuracy is used as the primary evaluation metric for SST-2. All experiments are run on a single NVIDIA T4 GPU in Google Colab. During training, we log loss, accuracy, learning rate, and other diagnostics to Weights and Biases, which we later use to analyze training dynamics and compare efficiency across models.

III. RESULTS

A. Overall Validation Performance

Table I and Figure 1 summarize the validation accuracy of the four pretrained models on SST-2. RoBERTa-base achieves the best performance with an accuracy of about 93.35%, followed by ALBERT-base-v2 and BERT-base-uncased with very similar accuracies around 92.3–92.4%. DistilBERT is noticeably worse at roughly 90.94%, reflecting the cost of model compression in this setting.

B. Training Efficiency

Figure 2 compares the total training time of the models on a single T4 GPU. DistilBERT is the fastest model, finishing training in roughly 1 700 seconds, while RoBERTa-base is the slowest at about 3 000 seconds. BERT-base and ALBERT-base-v2 lie in between. Combined with the accuracy results, this reveals a clear accuracy–efficiency trade-off: RoBERTa-base provides the best accuracy but at the highest cost, whereas DistilBERT offers the shortest training time but loses several points of accuracy.

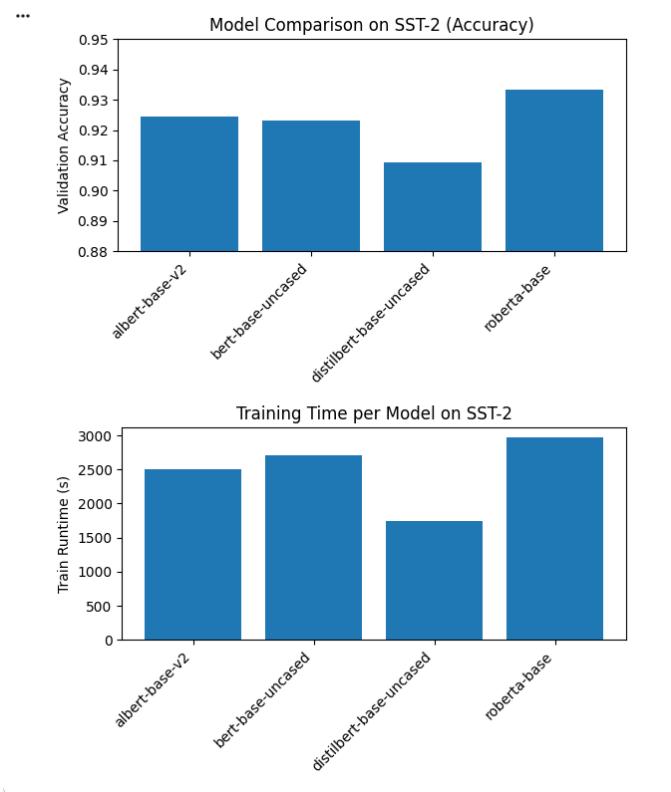


Fig. 2. Training time (in seconds) of each model on SST-2.



Fig. 3. Training loss versus global steps for the four models (exported from Weights and Biases).

C. Training Dynamics

To better understand how the models learn, we inspect the training and evaluation curves logged to Weights and Biases. Figure 3 shows the training loss as a function of global steps. All four models exhibit a smooth decrease and appear to converge after approximately 20k steps, with RoBERTa-base and BERT-base-uncased reaching slightly lower final loss values than ALBERT-base-v2 and DistilBERT.

Figure 4 presents the evaluation accuracy measured at several checkpoints during training. The curves confirm that RoBERTa-base consistently matches or outperforms the other models, while DistilBERT remains below the heavier encoders

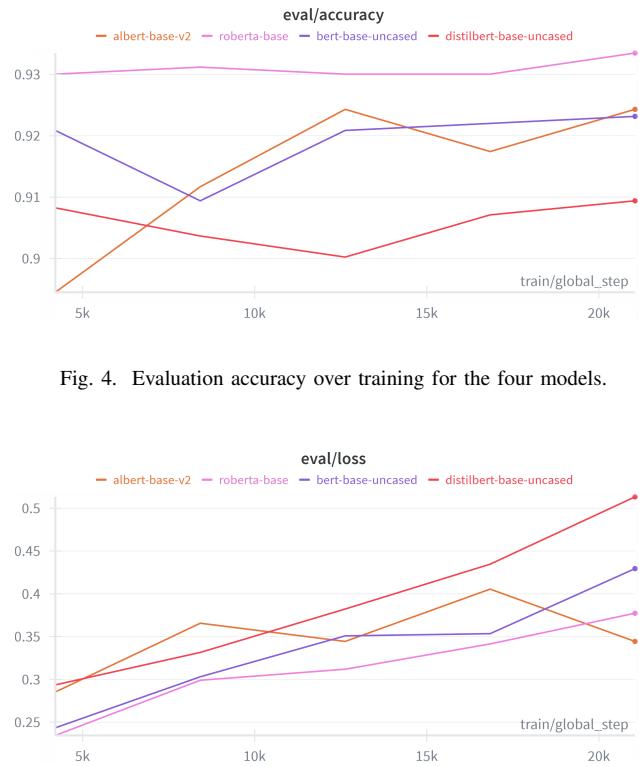


Fig. 4. Evaluation accuracy over training for the four models.

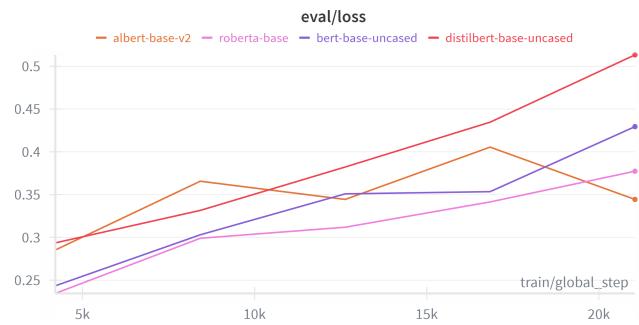


Fig. 5. Evaluation loss over training for the four models.

throughout training.

We also inspect the evaluation loss curves in Figure 5. The loss decreases rapidly during the first few thousand steps and then plateaus, again indicating that five epochs of fine-tuning are sufficient for this task.

D. Label Distribution

Finally, Figure 6 visualizes the label distribution in the SST-2 training set. The dataset is slightly imbalanced with more positive than negative examples, but the imbalance is not extreme.

IV. CONCLUSION

In this project we built a unified fine-tuning pipeline with the Hugging Face transformers and datasets libraries and used it to compare four widely adopted pretrained encoders—DistilBERT, BERT-base-uncased, ALBERT-base-v2, and RoBERTa-base—on the SST-2 sentiment classification task from the GLUE benchmark. By keeping the preprocessing, optimization hyperparameters, and evaluation protocol fixed across models, we were able to attribute differences in downstream performance primarily to the underlying architectures and pretraining strategies.

Our experiments show that RoBERTa-base achieves the best validation accuracy on SST-2 (about 93.3%), confirming its strong performance reported in the literature. ALBERT-base-v2 and BERT-base-uncased obtain very similar accuracies

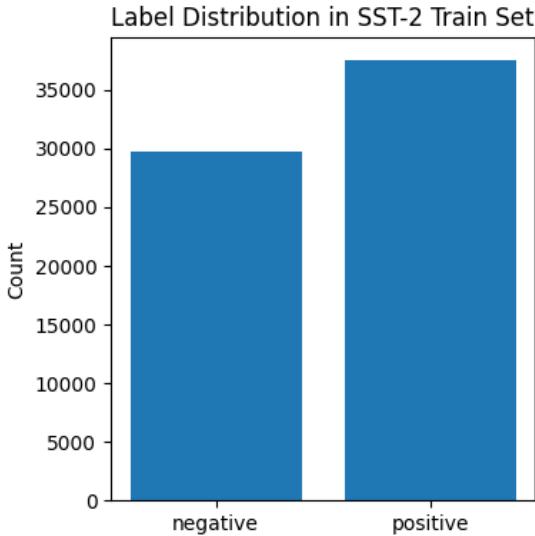


Fig. 6. Label distribution of the SST-2 training set.

around 92.3–92.4%, indicating that parameter sharing in ALBERT can reduce memory usage without significantly degrading performance on this task. DistilBERT, the smallest model considered, is several points worse at roughly 90.9% accuracy, but it trains substantially faster than the larger encoders. The training runtime measurements and W&B curves make the resulting accuracy–efficiency trade-off explicit: heavier models deliver modest but consistent gains in accuracy at the cost of longer training time and higher computational requirements.

These findings suggest the following practical guidance for sentiment classification on SST-2-like data. When accuracy is the primary objective and resources permit, RoBERTa-base is a strong default choice. When memory or time is constrained, ALBERT-base-v2 or BERT-base-uncased provide a good compromise between performance and efficiency. DistilBERT is attractive in scenarios where fast experimentation or deployment on limited hardware matters more than achieving the very best accuracy.

This study has several limitations. We only considered a single GLUE task and a single set of fine-tuning hyperparameters, and we evaluated training efficiency in terms of runtime on one T4 GPU rather than full inference latency or energy usage. Future work could extend this comparison to additional GLUE tasks, larger models (e.g., BERT-large and RoBERTa-large), and more aggressive hyperparameter tuning or learning-rate schedules. It would also be interesting to evaluate distilled and quantized variants under strict latency or memory budgets. Nevertheless, the present results already provide a concrete, reproducible case study of how different pretrained Transformer encoders behave on a standard sentiment classification benchmark.

V. REFERENCES

REFERENCES

- [1] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Proc. NAACL-HLT*, 2019.
- [2] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, “RoBERTa: A robustly optimized BERT pretraining approach,” *arXiv preprint arXiv:1907.11692*, 2019.
- [3] Z. Lan, M. Chen, S. Goodman, K. Gimpel, P. Sharma, and R. Soricut, “ALBERT: A lite BERT for self-supervised learning of language representations,” *arXiv preprint arXiv:1909.11942*, 2019.
- [4] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, “DistilBERT: A distilled version of BERT: smaller, faster, cheaper and lighter,” *arXiv preprint arXiv:1910.01108*, 2019.
- [5] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman, “GLUE: A multi-task benchmark and analysis platform for natural language understanding,” in *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, Brussels, Belgium, 2018, pp. 353–355.
- [6] R. Socher, A. Perelygin, J. Y. Wu, J. Chuang, C. D. Manning, A. Ng, and C. Potts, “Recursive deep models for semantic compositionality over a sentiment treebank,” in *Proc. EMNLP*, 2013.